

EEP 596 Final Project

Kuang-Ming Chen, Mike Huang

Outline

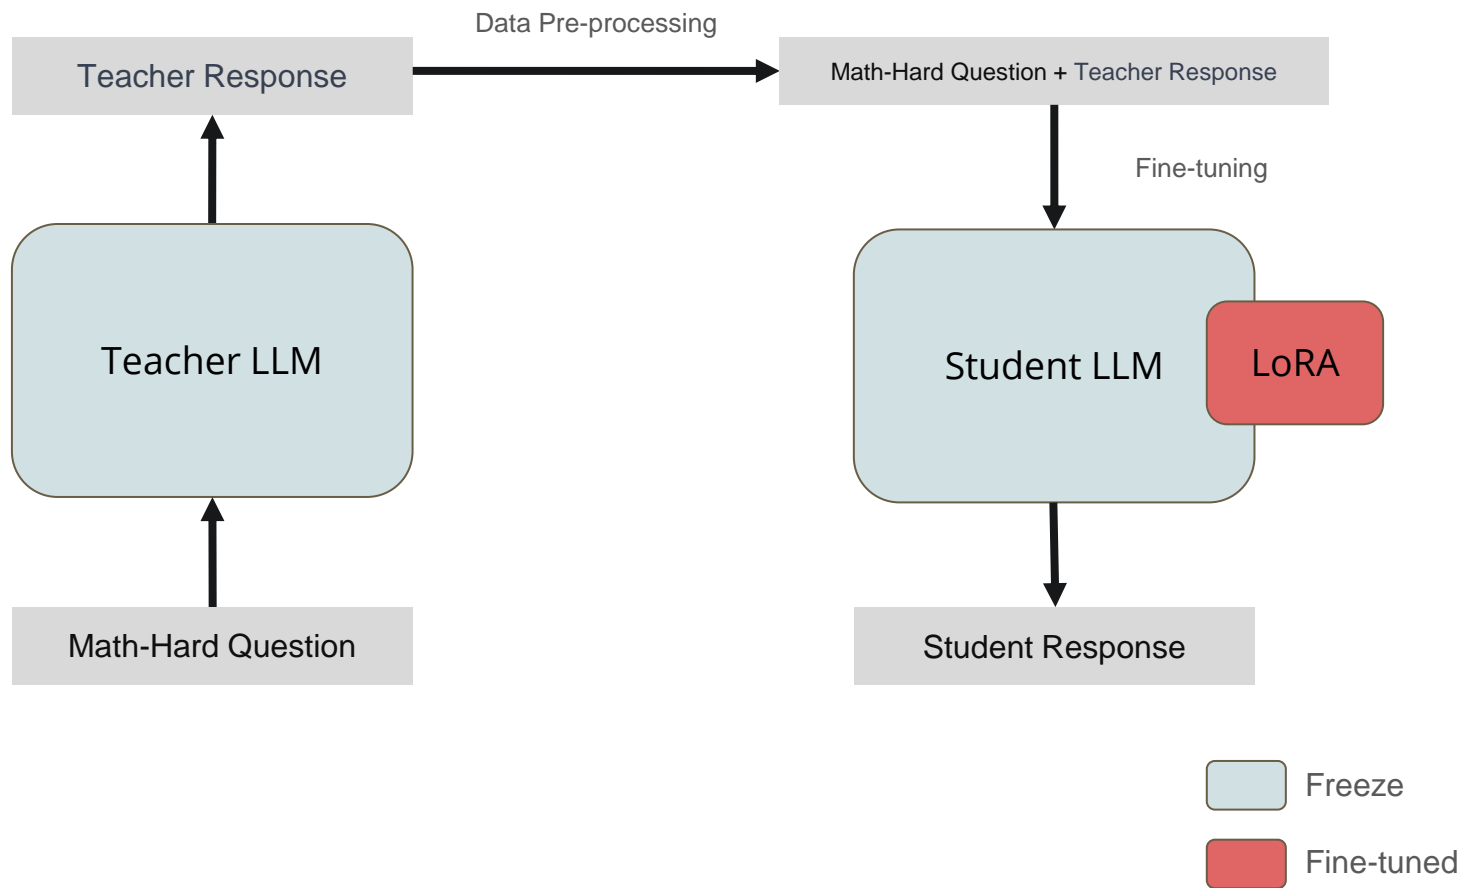
- Introduction and Motivation
- Dataset
- Experiment Settings
- Evaluation
- Experiment Results
- Reference

Introduction

- We are addressing a generation problem
- Goal: fine-tune LLM to solve difficult math questions by knowledge distillation
- Models: Llama-3.2-3B as teacher, Llama-3.2-1B as student
- Dataset: Math-Hard dataset
- Expectation: student model's performance will fall between student model with direct fine-tuning and teacher model's inference results





Motivation

- LLMs deploy on edge devices
- We need a smaller and high-performance model to solve real-world problems on these devices



Dataset

- Math-Hard dataset, 3.63k rows
- Use Problem and Solution columns in our project

problem	level	type	solution
string · lengths	string · classes	string · classes	string · lengths
			
174.31k	1 value	7 values	686.77k
What is the range of the function $y = \frac{x^2 + 3x + 2}{x+1}$? (Express your answer using interval...	Level 5	Algebra	We can factor the numerator to get $y = \frac{(x+2)(x+1)}{x+1}$. If we exclude the case whe...
Let $f(x) = \begin{cases} ax+3, & \text{if } x>2, \\ x-5 & \text{if } -2 \leq x \leq \dots \end{cases}$	Level 5	Algebra	For the piecewise function to be continu... cases must "meet" at $x=2$ and $x=-2$. For e...
Suppose a, b, c are positive numbers satisfying: $\begin{aligned} a^2/b &= 1, \\ b^2/c &= \dots \end{aligned}$	Level 5	Algebra	Notice that multiplying all three of the equations together tells us that $(a^2b^2c^2) \dots$
An infinite geometric series has a first term of $\$12$ and a second term of $\$4$. A second infinite...	Level 5	Algebra	Note that if the two series have con... ratios of a and b , respectively, then...
The entire graph of the function $f(x)$ is shown below (f is only defined when x is between -4 ...	Level 5	Algebra	First, we find all x such that $f(x) = 2$ by drawing the line $y = 2$ and finding the...
What is the y -coordinate of the point on the y -...	Level 5	Algebra	Because the point for which we are looki...

< Previous

1

2

3


...

24

Next >

Experiment Settings


- Use LoRA module to reduce computing resources and time
- Fix LLM weight, only fine-tuning on LoRA parameters



Experiment Settings

- Batch_size = 2
- Epoch = 1
- Learning_rate = $2e-4$
- Optimizer = AdamW
- Lr_scheduler = Linear
- LoRA_rank = 8

Experiment Settings



Training Loss (smoothing)

Evaluation

- MathQA dataset
- Multiple-choice questions, evaluate $P(\text{accuracy})$
- Expectation: student model's performance falls between baseline 2 and 3

Baseline 4	3B fine-tune on Math-Hard and inference
Baseline 3	3B direct inference
Baseline 2	1B fine-tune on Math-Hard and inference
Baseline 1	1B direct inference

Experiment Results

Model	Fine-tune	Baseline	Accuracy on MathQA
Llama-3.2-3B-Instruct (LoRA)	Math_Hard	Baseline 4	0.3534
Llama-3.2-3B-Instruct	-	Baseline 3	0.3451
Llama-3.2-1B-Instruct (LoRA)	Math_Hard_inf	Our Method	0.3333
Llama-3.2-1B-Instruct (LoRA)	Math_Hard	Baseline 2	0.3303
Llama-3.2-1B-Instruct	-	Baseline 1	0.3283

Thank you!

Reference

<https://huggingface.co/collections/meta-llama/llama-32-66f448ffc8c32f949b04c8cf>

<https://huggingface.co/datasets/lighteval/MATH-Hard>

<https://math-qa.github.io/>

<https://arxiv.org/pdf/2106.09685>

<https://github.com/EleutherAI/lm-evaluation-harness>